



Original coflow results: a high fleet average can hide short-request stalls

Fleet computation in [1000, 2000] ms



Short cohort per input: compute / active time



Right: $\text{sum}(\text{compute}) / \text{sum}(\text{completion} - \text{admission})$; 192 short requests per input, matched across policies (32K/128 + 32K/256). 5 ms setup.